



DSCI-560

Lecture 4: Data Science Dev Platforms/Languages Data Science Professional Practicum

Young Cho

Department of Electrical Engineering
University of Southern California

Reproducibility

- **Self-contained** configuration/program/data
 - Dependencies may have incompatibilities
 - May work in Python 3.6 but does not in Python 3.7.
 - Some libraries may require different versions
 - In one project, library A requires numpy 1.0, but numpy 2.0 for project B
- **Isolated** from machine dependencies
 - Certain SW functions might depend on HW configuration
- **Portable** from one machine to another
 - May need to move from one machine to another
 - May have run out of storage or memory
 - May need faster performance
 - May need to demonstrate on another machine

Development

- **Self-contained Environments**
 - Helps to manage working environments better
 - For each new project, create a new environment
- **Validation and Repeatability**
 - Immediate Validation and Working Starting Point
 - Reduces erroneous instruction interpretation
 - Reduces the possibility of using the wrong data
 - Presents exact details of the experiments

Virtual Environments

- Virtual Machines
 - Computer Emulation Software
 - HW/SW Operations are Emulated with SW
 - VMs interact HW via Light-weight SW Layer
 - Examples: VMWare, VirtualBox, QEMU, etc.
- Containers
 - Not SW Emulation
 - Application Layer Abstraction for SW and their dependencies
 - Containers are natively run on the HW with Light-weight Virtualization Layer
 - Examples: Docker, CoreOS rkt, Mesos, etc.

Virtual Machines

- Hypervisors
 - Light-weight software layer between VM and HW
 - Separates VMs from each other
 - Allocates physical processors, memory, and storage
 - Manages hosting operating systems
- Software on VM
 - Operating System with system binaries and libraries
 - All applications managed by the OS
- Most Modern Processors have VM Support
 - VM specific HW operations for acceleration


Containers

- Container Manager (i.e. Docker)
 - Virtualizes Operating System
 - Enables Application Layer Abstractions
- Container Operations
 - Host OS kernel and libraries are shared
 - Template of an environment is created within
 - Runs a snapshot of the system
 - Consistent behavior of an app.

Images for Containers

- Read-only template for Containers
- Sets up all software and libraries with specific environment configurations
- Enables the creation of various containers
- Can be used for reproducible data science results

VM and Containers




What's the Diff: VMs vs Containers

VMs	Containers
Heavyweight	Lightweight
Limited performance	Native performance
Each VM runs in its own OS	All containers share the SAME host OS
Hardware-level virtualization	OS virtualization
Startup time in minutes	Startup time in milliseconds
Allocates required memory	Requires less memory space
Fully isolated and hence more secure	Process-level isolation, possibly less secure

Performance comparison

- Getting applications from development to production involves creating disk images
- Fast image creation enables rapid testing and continuous deployment



Time (s)	VM (Vagrant)	Docker
MySQL	236	129
NodeJS	304	49

- Docker: 2-6x faster

Size comparison




Image size	VM	LXC	Docker
MySQL	1.68 GB	0.4 GB	112 KB
NodeJS	2.05 GB	0.6 GB	72 KB

Docker: 2-6x smaller

- VMs contain entire OS, and have larger images
- Docker stores only differences (application layer)



Container Images



Dockerfile

- Create images automatically using a build script: «Dockerfile»
- Can be versioned in a version control system like Git or SVN, along with all dependencies
- Docker Hub can automatically build images based on dockerfiles on Github

Reproducibility

- Dependencies
 - Less than 50% of software could be built or installed
 - Difficult to reproduce computational environment
 - Possible Solution: Docker Container
 - Software dependencies change, affecting results
 - Possible Solution: Software versioning
- Imprecise Documentation
 - Difficult to figure out how to install
 - Possible Solution Dockerfile records for dependencies
- Barriers to Adoption and Reuse
 - Difficulty to coordinate build tools/package managers
 - Persistent problem with this point

Development Tool: Python


- Started by Guido Van Rossum as a hobby
- Now, the Most Popular Language
 - Also means everyone already knows
- Open Source
- Versatile
- Lots of prebuilt packages
 - This is the KEY
- BUT, only used like a wrapper
 - Need to go beyond!!



Guido Van Rossum

The Most Popular Programming Languages

Share of the most popular programming languages in the world*



* Based on the PYPL-Index, an analysis of Google search trends
for programming language tutorials.



@StatistaCharts

Source: PYPL

statista

2025 IEEE Ranking based on Types of Developments

Language Rank	Types	Spectrum Ranking
1. Python	🌐💻	100.0
2. C	📱💻🔊	99.7
3. Java	🌐📱💻	99.5
4. C++	📱💻🔊	97.1
5. C#	🌐📱💻	87.7
6. R	💻	87.7
7. JavaScript	🌐📱	85.6
8. PHP	🌐	81.2
9. Go	🌐💻	75.1
10. Swift	📱💻	73.7

Python Today

- Development Today
 - Large and active scientific computing and data analysis community
- One of the most important languages for
 - Data science
 - Machine learning
 - General software development
- Many Prebuilt Packages

Two Modes

I. IPython

Python can be run interactively

Used extensively in research

2. Python scripts

What if we want to run more than a few lines of code?

Then we must write text files in .py

Installing Python

Windows:

- Download Python from
<http://www.python.org>
- Install Python.
- Run **Idle** from the Start Menu.

Mac OS X:

- Python is already installed.
- Open a terminal and run `python` or run **Idle** from Finder.


Linux:

- Chances are you already have Python installed. To check, run `python` from the terminal.
- If not, install from your distribution's package system.

Note: For step by step installation instructions, see the course web site.

Interpreted Languages

- **Interpreted**
 - Not compiled like Java
 - Code is written and then directly executed by an **interpreter**
 - Type commands into interpreter and see immediate results



Why Learn C?

- Python is high-level, easy, but slower.
- C is lower-level, compiled, and much faster.
- C gives control over memory and hardware.
- Many Python libraries (NumPy, TensorFlow) are written in C.



Key Differences

- Compiled vs Interpreted
 - Python: interpreted, dynamic
 - C: compiled, static typing
- Manual vs Automatic memory management
- Syntax differences: indentation vs braces
 { }

First Program

- Python:

```
print("Hello, World!")
```

- C:

```
#include <stdio.h>

int main() {
    printf("Hello, World!\n");
    return 0;
}
```

Variables & Types

Python:

- `x = 10`
- `y = 3.14`

C:

- `int x = 10; // 32-bit fixed point number`
- `float y = 3.14; // 32-bit floating point number`

Control Structures

Python:

- if $x > 0$:
- print("Positive")
- else:
- print("Non-positive")

C:

- if ($x > 0$) {
- printf("Positive\n");
- } else {
- printf("Non-positive\n");
- }

Loops

Python:

- `for i in range(5):`
- `print(i)`

C:

- `for (int i = 0; i < 5; i++)`
 {
- `printf("%d\n", i);`
- }

Functions

Python:

- `def add(a, b):`
- `return a + b`

C:

- `int add(int a, int b) {`
- `return a + b;`
- `}`

Arrays vs Lists

Python:

- arr = [1, 2, 3, 4]
- print(arr[2])

C:

- int arr[4] = {1, 2, 3, 4};
- printf("%d\n", arr[2]);

Strings

Python:

- `s = "Hello"`
- `print(len(s))`

C:

- `char s[] = "Hello";`
- `printf("%lu\n",`
`strlen(s));`

Memory Management

- Python: Automatic garbage collection
- C:

```
int *p = malloc(sizeof(int) * 5);  
free(p);
```

Pointers

- C Example:

```
int x = 10;  
int *p = &x;  
printf("%d\n", *p); // dereference
```

Structs vs Classes

Python:

- class Point:
- def __init__(self, x, y):
- self.x = x
- self.y = y

C:

- struct Point {
- int x;
- int y;
- };

Python ↔ C Interoperability

- Python can call C libraries with `ctypes` or `cffi`.
- Example
 - NumPy, TensorFlow use C under the hood.
- Use C for performance-critical code.

Python to C

- Python is high-level, dynamic, memory-safe.
 - C is low-level, explicit, and fast.
 - Knowing both = power + flexibility.
- Next Steps:
 - Write simple C programs.
 - Experiment with memory management.
 - Try calling C from Python.

Why CUDA?

- CPUs
 - A Few Powerful cores
 - Good for sequential tasks.
- GPUs
 - Thousands of lightweight cores
 - Ideal for parallel tasks.
- CUDA
 - NVIDIA's platform to program GPUs using C/C++.
- Used in AI, simulations, graphics, data science.

CUDA vs C

- C: Programs run on CPU.
- CUDA C: C + extensions for GPU programming.
- Functions can run on CPU (host) or GPU (device).
- Explicit memory management between CPU and GPU.

CUDA Programming Model

- Host (CPU) and Device (GPU).
- Functions:
 - `__host__` – runs on CPU
 - `__device__` – runs on GPU
 - `__global__` – GPU kernel callable from CPU
- Threads grouped into blocks and grids.

Hello CUDA

C Program:

- `printf("Hello from CPU\n");`

CUDA Program:

- `#include <stdio.h>`
- `__global__ void hello() {`
- `printf("Hello from GPU!\n");`
- `}`
- `int main() {`
- `hello<<<1, 1>>>();`
- `cudaDeviceSynchronize();`
- `return 0;`
- `}`

Thread Hierarchy

- Thread → smallest execution unit.
 - Block → group of threads.
 - Grid → group of blocks.
-
- CUDA IDs:

```
int tid = threadIdx.x;
int bid = blockIdx.x;
int gid = bid * blockDim.x + tid;
```

Example: Vector Addition

- CPU C Code:

```
for (int i = 0; i < N; i++) {  
    C[i] = A[i] + B[i];  
}
```

- CUDA Kernel:

```
__global__ void vecAdd(int *A, int *B, int *C,  
int N) {  
    int i = blockIdx.x * blockDim.x +  
threadIdx.x;  
    if (i < N) C[i] = A[i] + B[i];  
}
```

- Launch:

```
vecAdd<<<(N+255)/256, 256>>>(A, B, C, N);
```

Memory in CUDA

- Host Memory (CPU RAM) vs Device Memory (GPU VRAM).
- Explicit copies:

```
cudaMalloc( (void**) &d_A, size);  
cudaMemcpy(d_A, h_A,  
size, cudaMemcpyHostToDevice);
```

- Types:
 - Global (slow, big)
 - Shared (fast, per block)
 - Registers (fastest, per thread)

Shared Memory Example

```
__global__ void addShared(int *A, int *B, int *C)
{
    __shared__ int temp[256];
    int tid = threadIdx.x;
    temp[tid] = A[tid] + B[tid];
    __syncthreads();
    C[tid] = temp[tid];
}
```

- Shared memory allows fast collaboration within a block.

Performance Considerations

- - Use many threads to hide latency.
- - Coalesced memory access = faster.
- - Minimize CPU↔GPU transfers.
- - Use shared memory wisely.
- - Profile with nvprof or nsight.

CUDA vs Multithreading in C

- C with `pthreads/OpenMP`
 - Parallel on CPU only.
- CUDA: thousands of GPU threads.
- GPU excels at data-parallel problems.

CUDA Applications

- Deep Learning (PyTorch, TensorFlow)
- Image processing
- Physics simulations
- Financial modeling

C to CUDA

- CUDA extends C for GPU programming.
- Kernels, threads, blocks, memory are key concepts.
- Workflow: Allocate → Copy to GPU → Launch kernel → Copy back.
- Learn memory hierarchy, profiling, optimization.